

Transaction Management in DBMS

Transaction :- It is also a sequence of instructions / operations
(same like process, job task, thread..)

only diff. is
All the operations of the transaction will execute completely
at same time. (if some interrupt, then it won't stop there.)

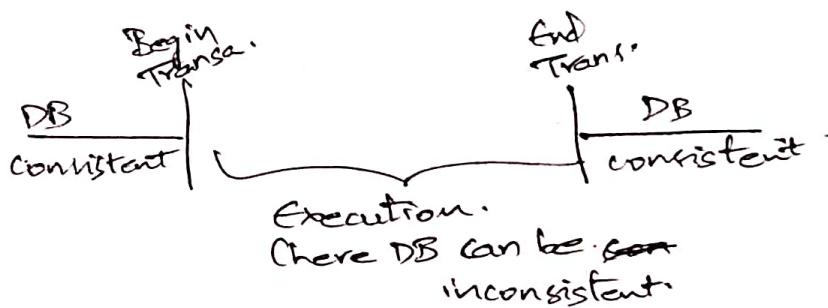
Transactions are atomic in nature (either complete totally or roll back)
if some.

→ A transaction is a single logical unit of work which access DB and possibly modify the DB.

ACID Properties:- Atomicity, consistency, Isolation, Durability.

Assumption:- initially database is consistent,
only transaction interacts with DB so it has to modify
the DB such that it transforms one consistent state to
another consistent state.

If it transforms to inconsistent state (No)



R(A)

A = A - 10

W(A)

R(B)

B + 10

W(B)

If an interrupt

happens here, then

it must roll back.

Atomicity :- either all the No partial execution of transaction is allowed

Consistency :- A transaction has to lead from one consistent state to another consistent state.

Ensure consistency.

Isolation:- Tatkal tickets, suppose transactions take lot of time, not possible > so simultaneously no. of transactions done.

Concurrency.

Race condition arises and leading to inconsistency

$T_1 \downarrow \downarrow T_2$ logical isolation:- Many ~~start~~ transactions simultaneously done.

But, its like only one is executing for them.

bit for T_1 , it thinks only its executing
Same for T_2 .

Intermediate results of individual transactions must not be available to each other.

Durability:- ~~that~~

If a transaction commits, whatever changes it does, persist forever.

Even if system fails, then can't roll back,

↓
How to roll back a committed Transaction
(not possible).

if R(A)

W(A)

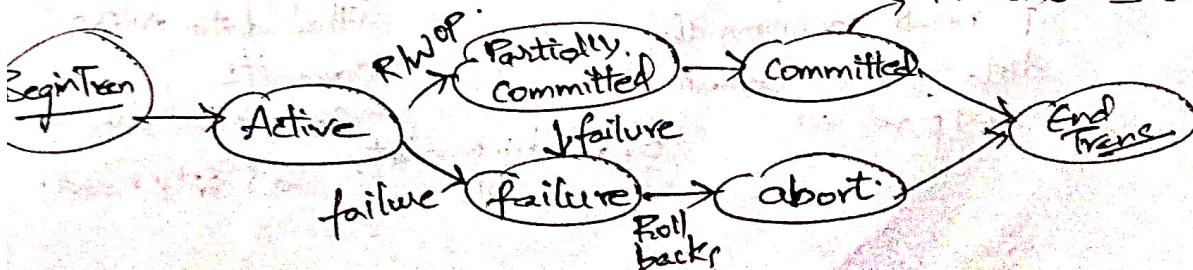
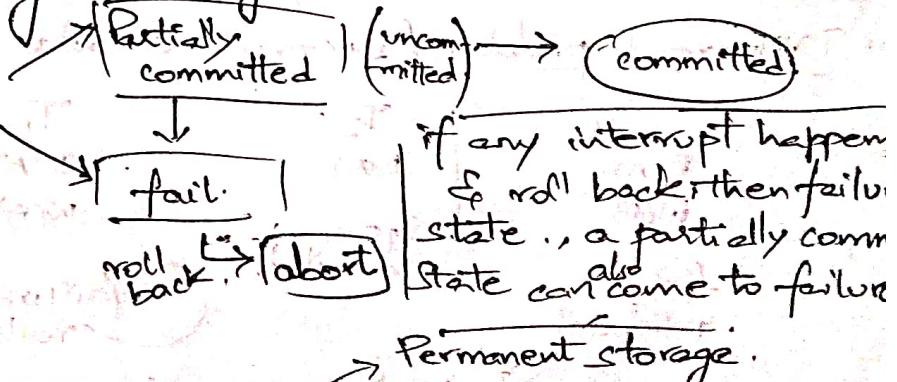
then we need unwrite, write another Transaction

Compensating Transaction

Transaction states:-

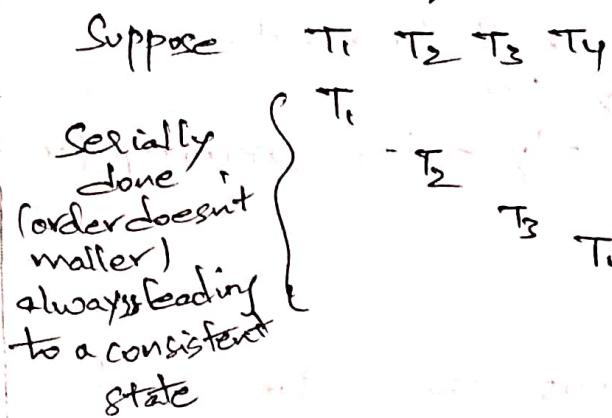
When a transaction begins, it goes into active state.

BeginTrx \rightarrow Active



Concurrency & problems associated with it

If, ATMs transactions are done serially, then no two persons can withdraw money at the same time. (not reliable)
So, concurrency.



Serial execution of Transaction:-

The execution is not overlapped.

Non-serial ex. of Trans:- (Concurrent ex. of Trans.)
• always allow overlapped Trans.

$T_1 \ T_2$

$T_1 \ T_3$

$T_2 \ T_4$

Advantage of concurrency:-

Drawback of Serial ex.:- no effectiveness, productivity loss.

- (1) Response Time \downarrow (because T_2 waits for T_1)
- (2) Avg. waiting Time \downarrow (because T_2 must wait for T_1 in serial)
- (3) Resource Utilization \uparrow
- (4) Efficiency \uparrow (because of all above, performance \uparrow)

But, concurrency may lead to inconsistent data.

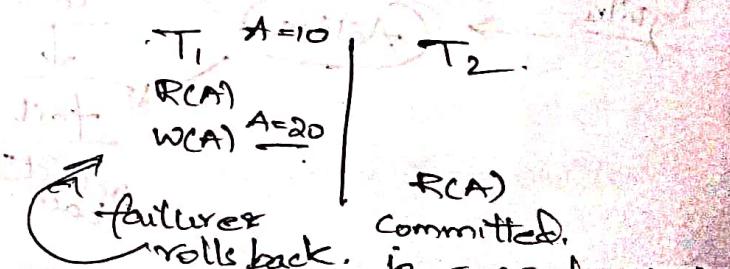
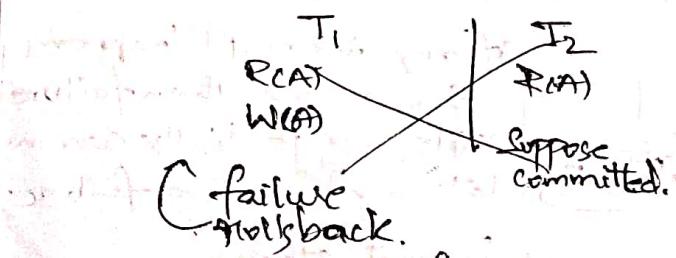
So, a controlled concurrency is allowed.

(controlled by ACID properties).

A transaction with ACID properties is allowed.

Problems:-

1) Dirty read problem / uncommitted data.

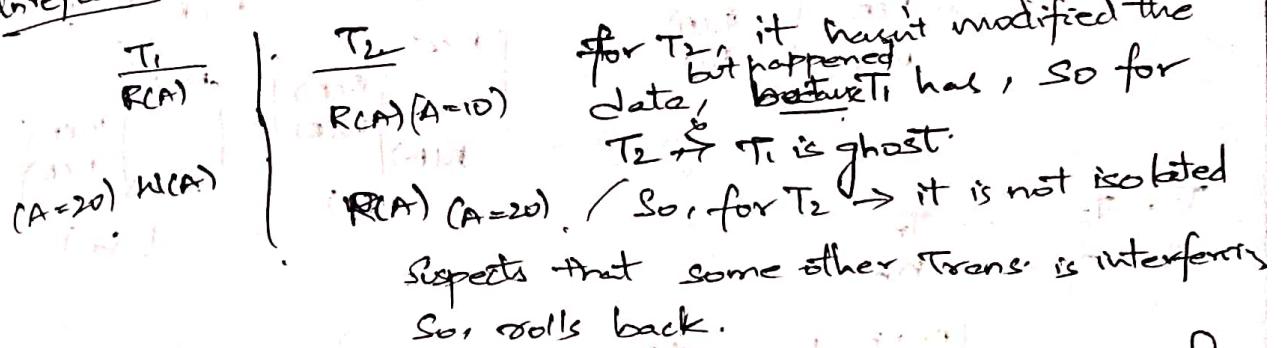


If read as committed data & then commit
no problem

So $A=20 \rightarrow$ wrong.
So reading dirty So Dirty Read.

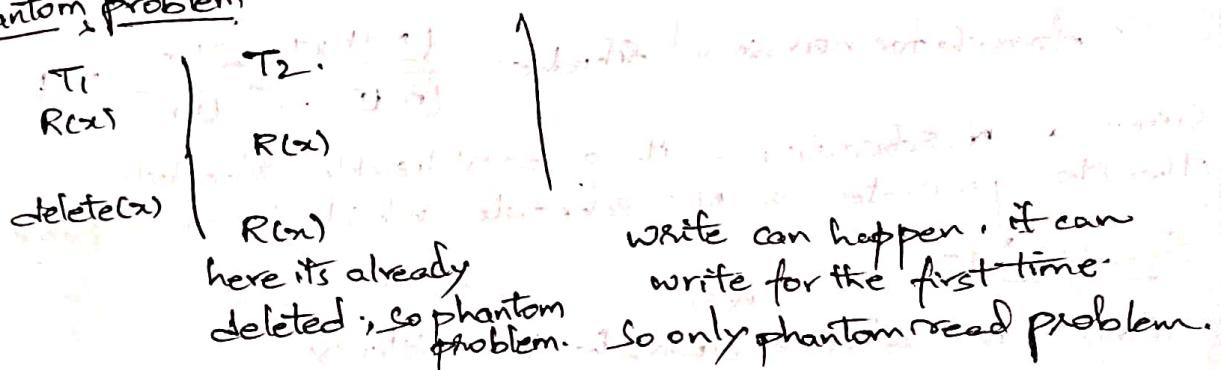
here read uncommittted data $A=20$ & commits.

- Solutions:-
- (1) Don't read uncommitted data.
 - (2) If read an uncommitted data, then don't commit, if the previous commits, then commit.
- (2) Unrepeatable read problem. (Due to isolation problems).

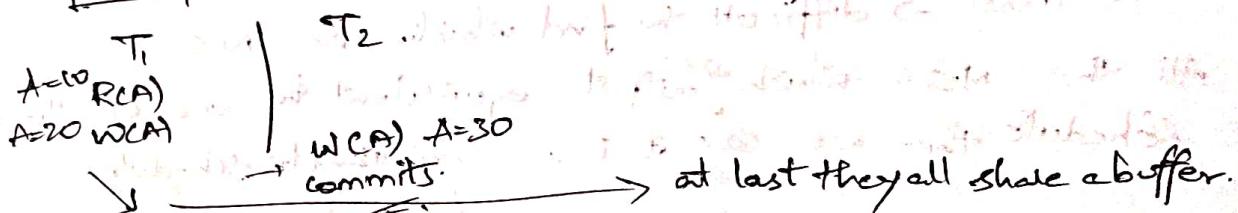


here in T_2 , read is repeated, so, problem arises
 So, don't repeat the read.

* Phantom problem



3) Lost Update problem:-



So, previous update $A=20$ is lost.

Schedules

dictates the order in which the transaction should execute.

T_1, T_2, \dots, T_n Transactions \rightarrow how they are scheduled.

$T_1, T_2, T_3 \rightarrow$ how many possible serial schedules.
 $3 \times 2 \rightarrow 6$ (3!)

Possible:

Suppose, n transactions then $n!$ serial scheduling.

Serial schedule Non-serial schedule

$T_1 \quad T_1 \quad T_2 \quad T_2 \quad T_3 \quad T_3$

$T_2 \quad T_2 \quad T_3 \quad T_3 \quad T_1 \quad T_1$

$T_3 \quad T_3 \quad T_2 \quad T_2 \quad T_1 \quad T_1$

$$S.S = n! \overline{s-T_1 - \dots - T_n}.$$

N.S.S - Many lead inconsistency.

$\tau_1, \tau_2, \dots, \tau_n$.

n_1, n_2, \dots, n_n the numbers to be arranged.

$$\Leftrightarrow \frac{n_1 + n_2 + \dots + n_n}{n_1 \cdot n_2 \cdot n_3 \cdot \dots \cdot n_n} - \ln$$

division to
remove inner
schedules

Subtract serial schedules.

\Rightarrow formula for non-serial schedules = $\frac{n_1 + n_2 + \dots + n_n}{n_1 \cdot n_2 \cdot \dots \cdot n_n} - 1$

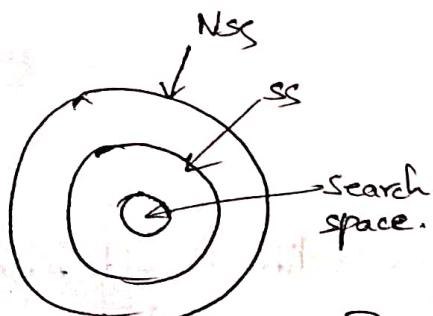
Given a NSSchedule, is it a serializable schedule.
How to generate a N.S. schedule which is serializable.

Serializability

Since, N.S. schedule may leads to inconsistency.

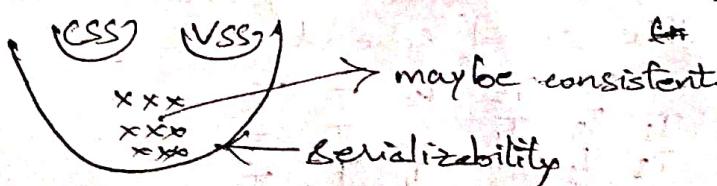
look trans \rightarrow difficult to find which are N.S.

All the N.S.S which are output equivalent to one of the serial schedule then we say it is a serializable schedule.



Both conflict serializable sch.
view serializable sche.

C.S.S & V.ISS always lead to consistency.
But if the given schedule is not C.S.S (or)
not V.ISS, it doesn't mean that it will
lead to inconsistency.



This problem is an NPC problem i.e. they have no algorithmic procedure but given an answer we can verify.

Conflict serializability:-
(Stricter)

Given a N.S.S, is it conflict S.S.
CSS = it has conflicting operations. 2)

View serializability.

(stricter) flexible.

(little strict)

Simultaneously done.

Read Read. \Rightarrow Non-conflict

Read write

simultaneously \Rightarrow Conflicting

W, R, \Rightarrow Conflicting

W, W \Rightarrow Conflicting

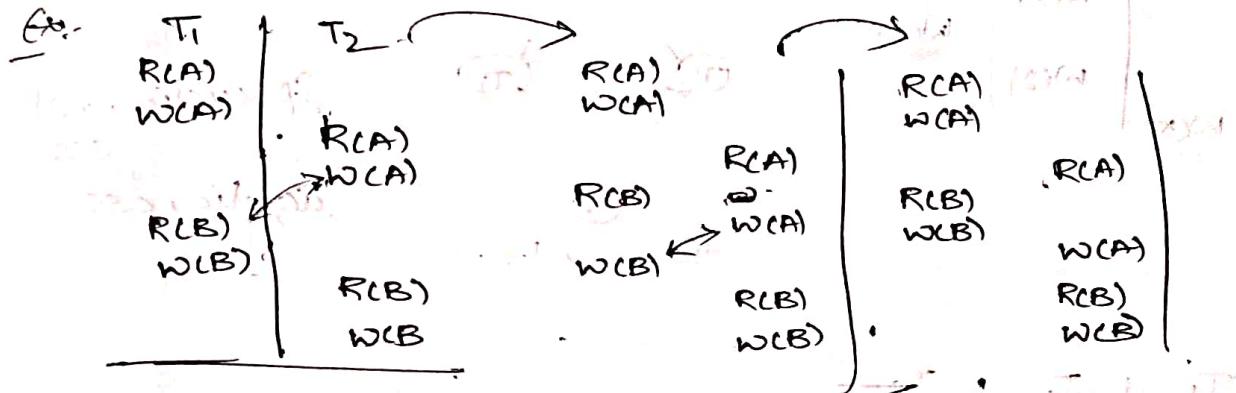
when done on something R(A) was conflicting.

R(A), W(B) -

Non conflicting

Conflict serializability.

If through series of swapping ~~operations~~ of non-conflicting operations, if you can convert N.S.S into one of the possible S.S then the schedule is C.S.S.

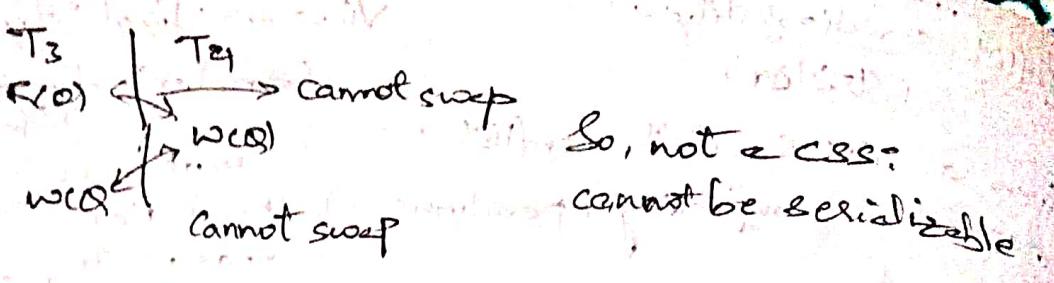


This is a serial schedule.

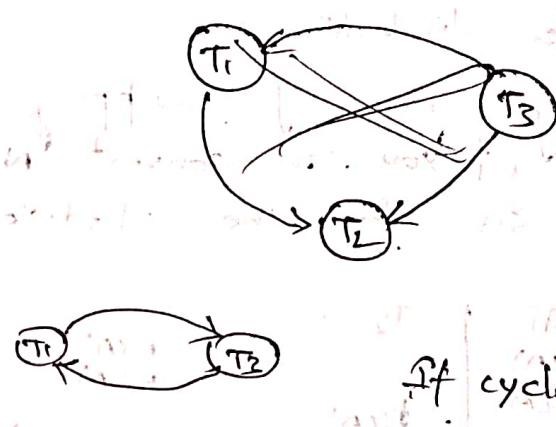
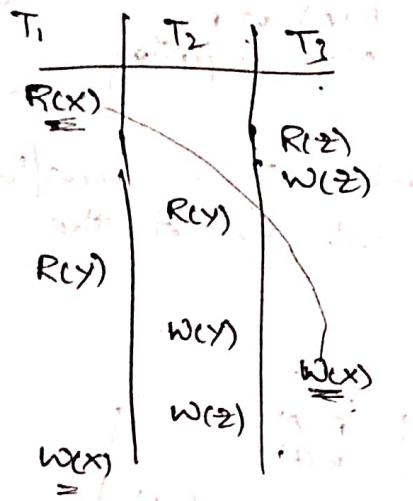


A NSS is conflict serializable iff it is conflict equivalent to any possible serial schedule.

Conflict equivalent:- If through series of swapping of non-conflicting operations if you can convert NSS into one of the possible S.S. S to S'.

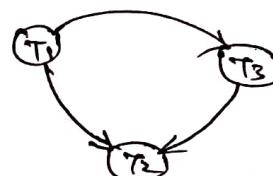
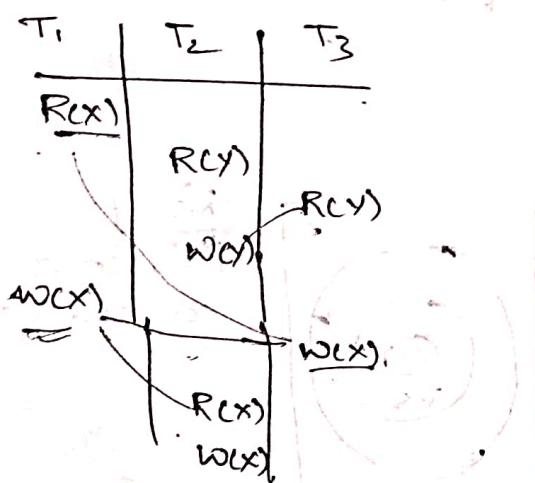


if no. of transactions goes, T_1, T_2, \dots, T_{100} , then difficult.
 Then draw a precedence graph.
 It is a directed graph.



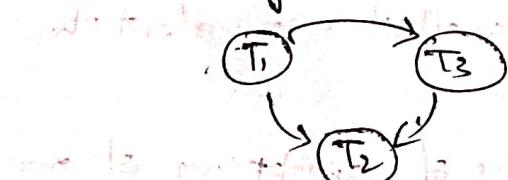
If cyclic, not CSS.
 acyclic, CSS.

② there cyclic.



acyclic
 So, CSS.

Precedence graph also tells order of serializability. \Rightarrow Topological sorting



T_1 depends on T_1, T_2
 so first T_1

& next $T_2 \rightarrow T_3$

so next, T_3

$\boxed{T_1 | T_3 | T_2}$

View Serializability :-

it is flexible, little stricter.

try to avoid VSS. / All CSs are VSS.

if it is CS, it will be VSS

Suppose S is not CS, (may be VSS)

→ and S does not have a blind write then, not VSS
→ but have a blind write, then cannot escape from VSS.

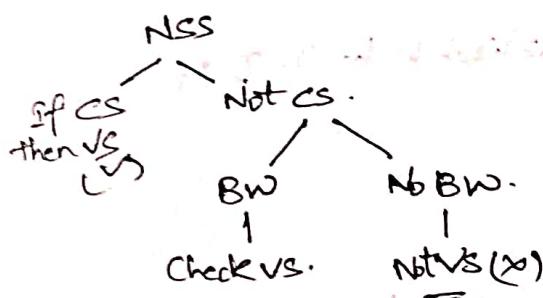
A schedule with blind write may not be VSS.

Blind write :-

if not CS, check for blind write

No blind write, cannot be VSS.

having a blind write is a necessary cond. for VSS but not sufficient.



A NSS is VSS iff it is view equivalent to one of the possible serial schedules.

View Equivalent :- check visually equivalent or not. to one of the SC. so if no. of Transactions increases, then becomes difficult.

Initial Read
Final Write
Update (or Intermediate)
read:

SIR	T ₁	T ₂
FW	T ₂	T ₂
Update	T ₁ → T ₂	T ₁ → T ₂

(NSS) (SC)
 S S'. S S'
 T₁ T₂ T₁ T₂
 X X

If T₁ is doing SIR in S', then S also T₁ must do SIR.

If T₂ is doing FW in S'; then in S also some Suppose T₁ reads & T₂ writes (T₁ → T₂) then in S also (T₁ → T₂)

Blindwriter

T_1	T_2
R(A)	
W(A)	

$W(A) \rightarrow$ in T_2 , without reading waiting, blindly writing.

Ex:-

T_1	T_2	T_3
R(A)		
	W(A)	
W(A)		

Blind writers

1 2 3

1 3 2

2 1 3

2 3 1

3 1 2

3 2 1

No intermediate reads

IR	T_1
FW	T_3
U	Ex.

So

$\underline{T_1} \underline{T_2} \underline{T_3}$ (Yes) view serializable to T_1, T_2, T_3 .

is it vs to 321 (No)

Ex:-

T_1	T_2	T_3	T_4
R(A)			
	R(A)		
		R(A)	
			R(A)
W(B)			
	W(B)		
		W(B)	
			W(B)

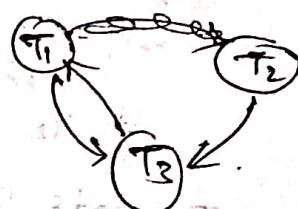
Not conflicting at all

So, ~~not ccs~~

So, surely yes.

Ex:-

T_1	T_2	T_3
R(A)		
	R(A)	
		W(A)
W(A)		



cyclic Not ccs.

So may be yes.

	A	
SR	T ₁	
FW	T ₁	

JMR. x.

(T₁) T₂ T₂ (T₁)

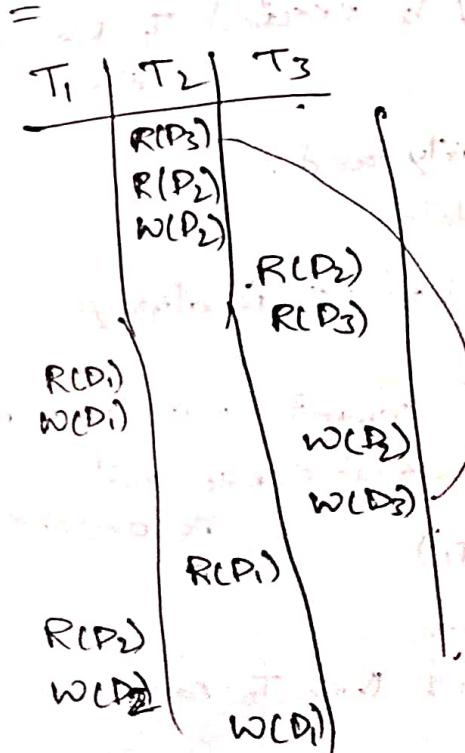
(b) not possible

So, not VS



	A	B
SR	T ₁	
FW	T ₂	T ₂

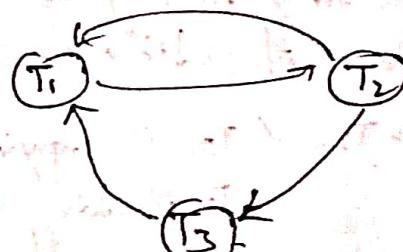
JMR



	P ₁	P ₂	P ₃
SR	T ₁	T ₂	T ₂
FW	T ₂	T ₂	T ₁

JMR

T₁ T₂ T₃ T₁ Not VS.

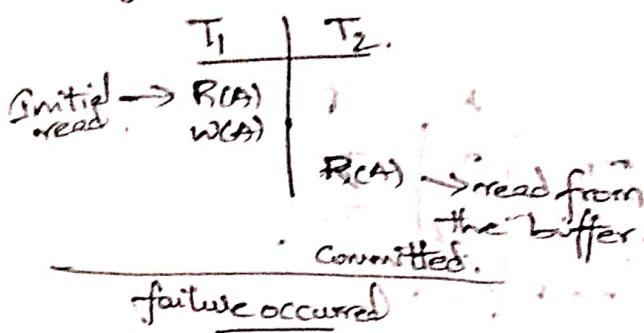


cyclic
Not cse.

The schedule is not serializable.

Recoverability:-

A schedule is recoverable, when all the transactions participating in the schedule are recoverable.



Recoverability is
almost (necessarily)
Property for a
schedule.

T₁ rolls back, but T₂ can't roll back.

R(A)=10 & here T₁ W(A)=20 → and T₂ read 20 elements
but T₁ rolled back so, now T₂ read 20 is wrong.
It caused inconsistency.

T ₁	T ₂
A=10 & R(A)	
A=20 & W(A)	

R(A) A=20
W(A) A=15
A=5
Committed.

& failure occurred

then, T₁ rolls back, & A been
-es '10' then T₂ is wrong.
If only T₂ executed T₂ R(A)=10

T₂ Avoid dirty read.
always read committed data.

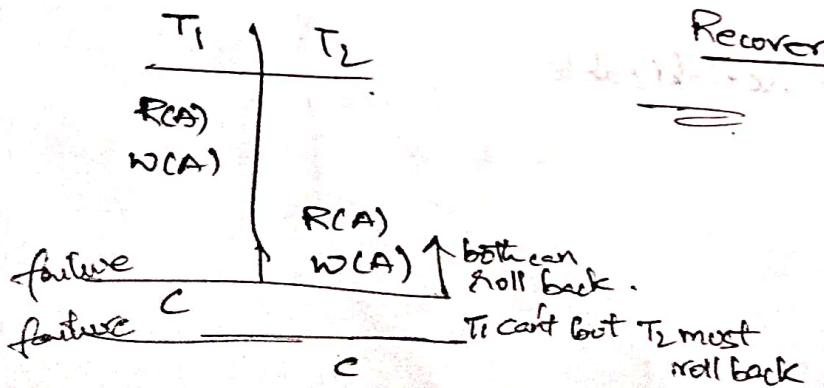
If there is no dirty read, the schedule is always
recoverable.

If there is a dirty read, the order of commit must correspond to order of dirty read, then recoverable else not.

T₁ → T₂ (T₂ is copying from T₁)

∴ T₂ is doing dirty read

then T₁ must commit first then T₂ can commit.



Recoverable:- It can roll back

T₁ → T₂ → T₃.

Commitment order:-

T₁ T₂ T₂
C₁ C₂ C₃.

Cascadelessness

$T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_4 \rightarrow T_5$.

\leftarrow
cascading rollbacks

if T_1 failure, then all ~~next~~ T_2, T_3, T_4, T_5 must roll back.
One causes many to fail.

cascading abort:- The abort of one transaction forces the abort of another transaction to prevent the second transaction from reading invalid (uncommitted) data.

If no dirty read, then cascadelessness(\checkmark)

If there is a dirty read, \rightarrow always cascading rollbacks.

T_1	T_2	T_3
R(A)		
w(A)		
C		

T_1	T_2	T_3
	R(A)	
	w(A)	
	C	R(A)

always read committed data,
so reading updated copy so,
no need to roll back.

All cascading cascadelessness schedules are recoverable
but not vice versa.



Strict schedule:- close to a serial schedule.

No dirty read, No dirty write.

always work on committed data.

T_1	T_2
A=10	R(A)
A=20	w(A)

w(A)
(A=5) A=15.

If there is a dirty write,
always a change of last
update problem, here
 T_1 lost its update.

Committed.
if thinks.
it is writing 20
but (no) 15.
So this is
dirty write.

T_1	T_2
R(A)	
w(A)	
C	

here no
dirty write
because T_1
committed.

T_1	T_2
R(A)	
w(A)	
C	

T_1	T_2
R(A)	R(B)
w(A)	w(B)
C	C

T_1	T_2
$R(A)$	
	$R(B)$
$W(A)$	
C	
	$W(B)$
	$R(A)$

This is a
strict
schedule

T_1	T_2
$R(A)$	
$W(A)$	
C	
	$W(A)$
	$R(A)$
	C

No dirty read

So, cascadelessness

It is not a strict schedule.

→ This is a dirty write (lost update problem)
(writing uncommitted data)

So first commit, then write.

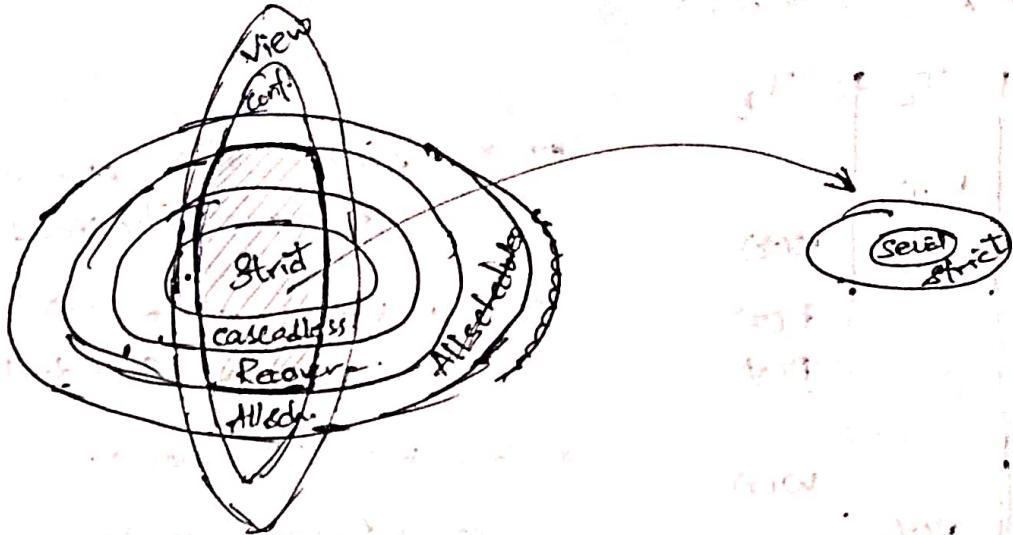
All serial schedules are strict schedules (but not vice versa).

T_1	T_2	T_3
$R_1(X)$		
$R_1(Z)$	$R_2(X)$	
$R_1(Z)$		
		$R_3(Z)$
$W_1(X)$		
C_1		
		$R_3(Y)$
	$R_2(Y)$	
	$W_2(Z)$	
	$W_2(Y)$	
		C_3
		C_2

It is non-serial,
but strict schedule
(No dirty reads & no
dirty write)

Relation among all type of schedules

We cannot work on irrecoverable & view scheduler.



(1) S: R₂(B) W₂(A) R₁(A)
R₃(A) W₁(B) W₂(B) W₃(B)

If is not conflicting serializable.
But there is blind write.

T ₁	T ₂	T ₃
R(A)	R(B)	R(A)
W(B)	W(B)	W(B)

	A	B	C
IR	No initial record		
FW		T ₂	T ₃
I/M.R.	X		

→ first writing so, T₁ is not initial read.

	A	B	C
IR	T ₁	T ₃	T ₂
FW	T ₁	T ₂	T ₃
I/M.R.	-	-	-

(2). T₁: R(A)
T₂: R(A)
T₃: R(B)
W(A)
R(C)
R(B)
W(B)
W(C)

Not viewserializable.

T ₁ T ₂ T ₂	T ₂ T ₃ T ₁	T ₃ T ₁ T ₂
T ₁ T ₃ T ₂	T ₂ T ₁ T ₃	T ₃ T ₂ T ₁

Recoverability examples:

(3) S: $R_1(x), R_2(x), R_1(z), R_3(z), R_2(y), W_1(x), W_3(y), R_2(y), W_1(y)$,
 $W_2(y), C_1, C_2, C_3$:

T_1	T_2	T_3
$R(x)$	$R(x)$	
$R(z)$		$R(z)$
		$R(x)$
		$R(y)$
$W(x)$		$W(y)$
		$R(y)$
		$W(y)$
		$W(y)$
C_1	C_2	C_3

on x , after $W_1(x)$, no dead
so, no dirty read

on y , $W_3(y) \rightarrow R_2(y)$
so, dirty read.

on z , no dirty read.

commitment order,

so, recoverable

→ Cascading rollback (\because dirty read)

Not a strict schedule (\because dirty read
with a dirty & dirty write).

(4) S: $R_1(x), W_1(x), R_2(x), R_1(y), R_2(y), W_2(x), W_1(y), C_1, C_2$:

T_1	T_2	
$R(x)$		
$W(x)$		
$R(y)$	$R(x) \rightarrow$	
	$R(y)$	
	$W(x)$	
$W(y)$		
C_1		
	C_2	

on x , dirty read.

T_2 is reading from T_1 .

so, T_1 commit first (✓).

so, recoverable.

Cascading rollback (\because dirty read)

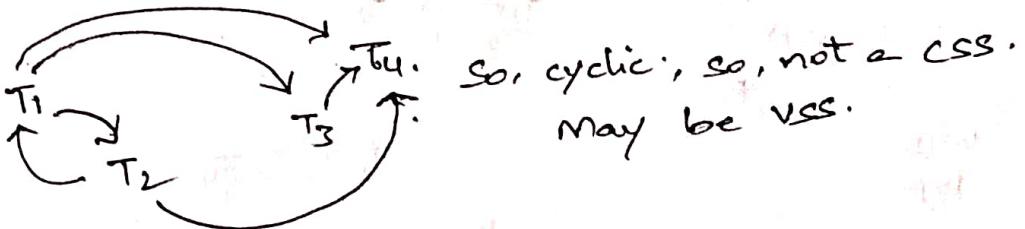
Not a strict schedule (\because dirty read).

(5) $R_2(A), R_1(A), W_1(C), R_3(C), W_1(B), R_4(B), W_3(A), R_4(C),$
 $W_2(D), R_2(B), W_4(A), W_4(B)$

T_1	T_2	T_3	T_4
	$R(A)$		
$R(A)$			
$W(C)$			
$W(C)$			
$W(B)$			
	$R(C)$		
		$R(C)$	
		$R(B)$	
			$R(C)$
			$R(C)$
			$R(C)$

T_1	T_2	T_3	T_4
	$W(D)$		
	$R(B)$		
			$W(A)$
			$W(B)$

on C, $w_1(c) \rightarrow R_A(c) \rightarrow$ There is a dirty read &
 don't have a commitment order, so can't conclude
 Since there is a dirty read, so cascading roll back.
 Since there is dirty read & dirty write, so, not a strict
 schedule.



	A	B	C	D
IR	T1	T2	T3	-
FW	T4	T4	-	T2
SIMR	-	$T_1 \rightarrow T_4$	$T_1 \rightarrow T_3$	-
	$T_1 \rightarrow T_2$	$T_1 \rightarrow T_4$	$T_1 \rightarrow T_3$	

	T1	T2	T3
R(A)	R(A)		
	R(A)	R(B)	
W(A)	R(C)	R(B)	
	R(C)	R(B)	
		W(B)	
			WCC

	A	B	C
IR	T_1, T_2	T_3, T_2	T_2
FW	T_1	T_2	T_3
SIMR	$T_2 \rightarrow T_1$	$T_3 \rightarrow T_2$	-
	$T_2 \rightarrow T_1$	$T_3 \rightarrow T_2$	-
		$T_3 \rightarrow T_2$	$T_2 \rightarrow T_3$
		$T_3 \rightarrow T_2$	$T_2 \rightarrow T_3$

to make T_3 as initial read
 so, $w_2(B)$ comes first

$T_3 \rightarrow T_2, T_2 \rightarrow T_3$
 contradicting each other

	T1	T2	T3
	R(B)		
R(A)		R(A)	
	W(A)		
WCB	R(A)		
		R(C)	
		W(C)	
			WCC

	A	B
IR	x	T_2
FW	T_2	T_3
SIMR	$T_2 \rightarrow T_1$	x
	$T_2 \rightarrow T_3$	x
	$T_2 \rightarrow T_1$	
	$T_2 \rightarrow T_3$	
		$T_2 \rightarrow T_3$
		$T_2 \rightarrow T_3 \rightarrow T_1$
		$T_2 \rightarrow T_3 \rightarrow T_1$

✓ V.S.S)

T ₁	T ₂	T ₃	T ₄
R(A)			
WCC)			RCC)
W(B)			R(B)
		W(CA)	R(CC)
W(CD)			
R(B)			
			W(CA)
			BU(E)

	A	B	C	D
IRR	T1 T2	0	0	0
FW	T4	T4	T1	T2
SIMR	0	T1 → T3 T1 → T2	T1 → T3 T1 → T4	0

T₁ T₂ T₃ T₄

It is view Serializable.

T_1	T_2	T_3	T_4	T_5
$R(A)$		$R(D)$		
$w(B)$	$R(B)$	$w(C)$	$R(B)$	
				$R(C)$
			$w(E)$	$R(E)$
				$R(A)$

	A	B	C	D	E
IIR	T ₁ T ₅	x	x	T ₃	x
FW	x	T ₃	T ₂	x	T ₄
IIMR.	x	T ₁ → T ₂ T ₃ → T ₄	T ₂ → T ₅	x	T ₄ → T ₅

$T_1 \ T_2 \ T_3 \ T_4 \ T_5$.

It is view Serializable.

T_1	T_2	T_3	T_4
$R(A)$			
$w(C)$		$R(c)$	
$w(B)$		$R(E)$	$R(CB)$
		$w(A)$	$R(CC)$
	$w(D)$		
$R(B)$			
		tot	$w(CA)$
			$w(CB)$

	A	B	C	D
IR	$T_1 T_2$	\cancel{x}	\cancel{x}	x
FW	T_4	$T_4.$	T_1	T_2
IMR.	\cancel{x}	$T_1 \rightarrow T_4$ $T_1 \rightarrow T_2$	$T_1 \rightarrow T_3$ $T_1 \rightarrow T_4$	x

\checkmark T₁ T₂ T₃ T₄

$\checkmark T_1 T_3 T_2 T_4$

Protocols that ensure to generate schedule which satisfies this property specifically CS.

$T_1, T_2, \dots, T_n, n=10000, \dots$

↓
Protocol

It generates a NSS following the properties - CS; VS; R.

Protocols

deadlock.

P₁
CS
R

P₂
CS
R

prefer P₂.

deadlock

A lock which is dead, permanent lock, never gets unlocked.

handles 2 processes
quad core, dual core, octa core.

Program → Passive entity.
Paint → active entity.

Every process needs some resources.

Ex: P₁ needs a resource A, which is already given to P₂

P₂ " " " "

P₁ should release A & then P₂ works & release B.



This is deadlock.

(Ctrl + Alt + Del) to kill all the processes.

deadlock → Infinite Wait.



Kill P₃

Starvation
(it cannot complete)

starves for CPU.

preempt
(forceful)

for deadlock there must be Mutual exclusion.

(all are necessary & sufficient)

Circular Wait.

Wait & Hold

No Preemption

kill lower priority processes, not high priority processes.

Protocols:-

(1) Timestamping Protocol (2) Thomas Write Rule.

3) Locking protocols 2PL (Two phase locking)

Validation Protocols.

Basic 2PL
Conservative 2PL
Rigorous 2PL
Strict 2PL.