

Scheduling Algorithms

(First-Come, First-Served Scheduling)

- By far the simplest CPU-scheduling algorithm.
- The process that requests the CPU first is allocated the CPU first.
- The implementation of the FCFS policy is easily managed with a **FIFO queue**.



- When a process enters the ready queue, its PCB is linked onto the tail of the queue.
- When the CPU is free, it is allocated to the process at the head of the queue.
- The running process is then removed from the queue.

Waiting time

The CPU scheduling algorithm does not affect the amount of time during which a process executes or does I/O; it affects only the amount of time that a process spends waiting in the ready queue. Waiting time is the sum of the periods spent waiting in the ready queue.

Response time

In an interactive system, turnaround time may not be the best criterion. Often, a process can produce some output fairly early and can continue computing new results while previous results are being output to the user. Thus, another measure is the time from the submission of a request until the first response is produced. This measure, called response time, is the time it takes to start responding, not the time it takes to output the response. The turnaround time is generally limited by the speed of the output device.

Scheduling Criteria

CPU utilization



Throughput

Turnaround time

Waiting time

Response time



Consider the following set of processes that arrive at time 0

Process	Burst Time (ms)
P1	24
P2	3
P3	3

If the processes arrive in the order P1, P2, P3, and are served in FCFS order, we get the result shown in the following Gantt chart:



Waiting Time for P1 = 0 ms

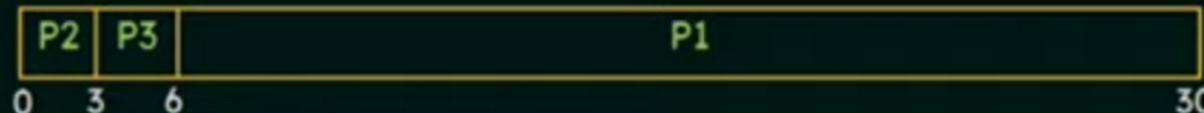
Waiting Time for P2 = 24 ms

Waiting Time for P3 = 27 ms

$$\text{Average Waiting Time} = (0 + 24 + 27)/3 = 17 \text{ ms}$$



If the processes arrive in the order P₂, P₃, P₁, however the result will be shown in the following Gantt chart:



Waiting Time for P₁ = 6 ms

Waiting Time for P₂ = 0 ms

Waiting Time for P₃ = 3 ms

$$\text{Average Waiting Time} = (6 + 0 + 3)/3 = 3 \text{ ms}$$



If the processes arrive in the order P₂, P₃, P₁, however the result will be shown in the following Gantt chart:



Waiting Time for P₁ = 6 ms

Waiting Time for P₂ = 0 ms

Waiting Time for P₃ = 3 ms

Average Waiting Time = $(6 + 0 + 3)/3 = 3$ ms

This reduction is substantial. Thus, the average waiting time under an FCFS policy is generally not minimal and may vary substantially if the process's CPU burst times vary greatly.



This reduction is substantial. Thus, the average waiting time under an FCFS policy is generally not minimal and may vary substantially if the process's CPU burst times vary greatly.

The FCFS scheduling algorithm is nonpreemptive

- Once the CPU has been allocated to a process, that process keeps the CPU until it releases the CPU, either by terminating or by requesting I/O.
- The FCFS algorithm is thus particularly troublesome for time-sharing systems, where it is important that each user get a share of the CPU at regular intervals.
- It would be disastrous to allow one process to keep the CPU for an extended period.



Scheduling Algorithms (Shortest-Job-First Scheduling)

- This algorithm associates with each process the length of the process's next CPU burst.
- When the CPU is available, it is assigned to the process that has the smallest next CPU burst.
- If the next CPU bursts of two processes are the same, FCFS scheduling is used to break the tie.

The SJF algorithm can be either preemptive or nonpreemptive

A more appropriate term for this scheduling method would be the

→ **Shortest-Next-CPU-Burst Algorithm**

because scheduling depends on the length of the next CPU burst of a process, rather than its total length.



Example of SJF Scheduling (Non-Premptive)

Consider the following set of processes, with the length of the CPU burst given in milliseconds:

Process ID	Burst Time
P1	6
P2	8
P3	7
P4	3

Waiting Time for P1 = 3 ms

Waiting Time for P2 = 16 ms

Waiting Time for P3 = 9 ms

Waiting Time for P4 = 0 ms

Gantt Chart:



Average Waiting Time

$$= (3 + 16 + 9 + 0)/4 = 7 \text{ ms}$$

By comparison, if we were using the FCFS scheduling scheme, the average waiting time would be 10.25 milliseconds.



Example of SJF Scheduling (Preemptive)

Consider the following four processes, with the length of the CPU burst given in milliseconds and the processes arrive at the ready queue at the times shown:

Process ID	Arrival Time	Burst Time
P1	0	8 - 7
P2	1	4
P3	2	9
P4	3	5

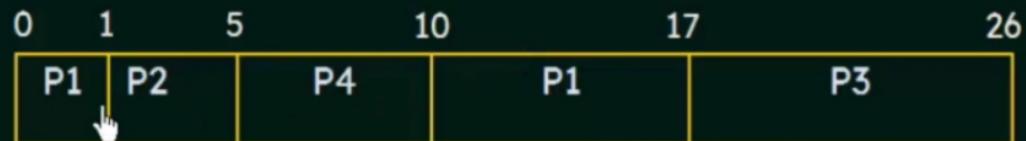
$$\text{Waiting Time for P1} = (10 - 1 - 0) = 9 \text{ ms}$$

$$\text{Waiting Time for P2} = (1 - 0 - 1) = 0 \text{ ms}$$

$$\text{Waiting Time for P3} = (17 - 0 - 2) = 15 \text{ ms}$$

$$\text{Waiting Time for P4} = (5 - 0 - 3) = 2 \text{ ms}$$

Gantt Chart:



Average Waiting Time

$$= (9 + 0 + 15 + 2)/4 = 6.5 \text{ ms}$$

Waiting Time = Total waiting Time - No. of milliseconds Process executed - Arrival Time



Problems with SJF Scheduling:

- The real difficulty with the SJF algorithm is knowing the length of the next CPU request.
- Although the SJF algorithm is optimal, it cannot be implemented at the level of short-term CPU scheduling.
- There is no way to know the length of the next CPU burst.

One approach is:

- To try to approximate SJF scheduling.
- We may not know the length of the next CPU burst, but we may be able to predict its value.
- We expect that the next CPU burst will be similar in length to the previous ones.
- Thus, by computing an approximation of the length of the next CPU burst, we can pick the process with the shortest predicted CPU burst.



Scheduling Algorithms (Priority Scheduling)

- A priority is associated with each process, and the CPU is allocated to the process with the highest priority.
- Equal-priority processes are scheduled in FCFS order.
- An SJF algorithm is simply a priority algorithm where the priority is the inverse of the (predicted) next CPU burst.
The larger the CPU burst, the lower the priority, and vice versa.

Priority scheduling can be either **preemptive** or **nonpreemptive**.

A **preemptive priority** scheduling algorithm will preempt the CPU if the priority of the newly arrived process is higher than the priority of the currently running process.

A **nonpreemptive priority** scheduling algorithm will simply put the new process at the head of the ready queue.

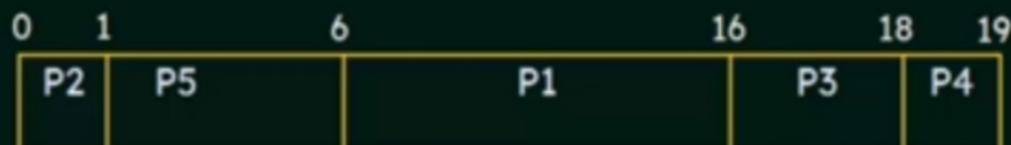


Consider the following set of processes, assumed to have arrived at time 0, in the order P1, P2, P3, P4, P5, with the length of the CPU burst given in milliseconds:

Process ID	Burst Time	Priority
P1	10	3
P2	1	1
P3	2	4
P4	1	5
P5	5	2

Waiting Time for P1 = 6 ms
Waiting Time for P2 = 0 ms
Waiting Time for P3 = 16 ms
Waiting Time for P4 = 18 ms
Waiting Time for P5 = 1 ms

Using Priority Scheduling, we would schedule these processes according to the following Gantt Chart:



Average Waiting Time

$$\begin{aligned} &= (6 + 0 + 16 + 18 + 1) / 5 \\ &= 41 / 5 \text{ ms} \\ &= 8.2 \text{ ms} \end{aligned}$$



Problem with Priority Scheduling

- A major problem with priority scheduling algorithms is indefinite blocking, or **starvation**.
- A process that is ready to run but waiting for the CPU can be considered blocked.
- A priority scheduling algorithm can leave some low priority processes waiting indefinitely.
- In a heavily loaded computer system, a steady stream of higher-priority processes can prevent a low-priority process from ever getting the CPU.

Solution to the Problem

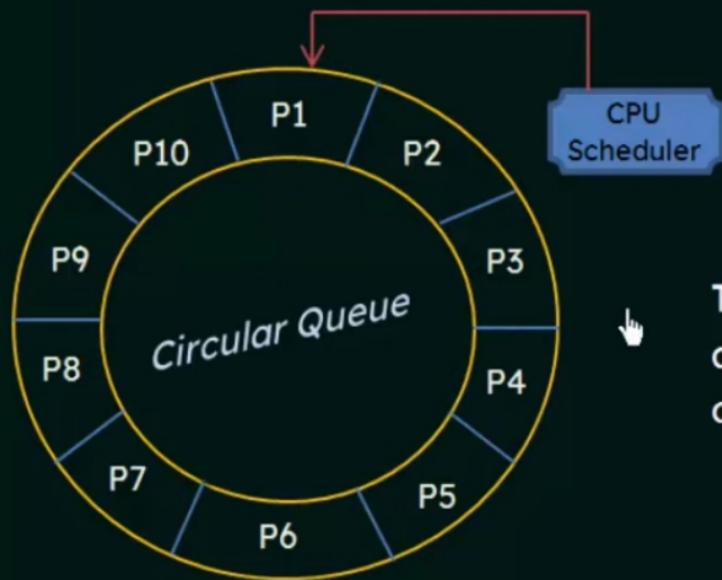
- A solution to the problem of indefinite blockage of low-priority processes is **aging**.
- Aging is a technique of gradually increasing the priority of processes that wait in the system for a long time.
- For example,
If priorities range from 127 (low) to 0 (high), we could increase the priority of a waiting process by 1 every 15 minutes.
- Eventually, even a process with an initial priority of 127 would have the highest priority in the system and would be executed.



Scheduling Algorithms

(Round-Robin Scheduling)

- The round-robin (RR) scheduling algorithm is designed especially for timesharing systems.
- It is similar to FCFS scheduling, but **preemption** is added to switch between processes.
- A small unit of time, called a time quantum or time slice, is defined (generally from 10 to 100 milliseconds)



- The ready queue is treated as a circular queue.

The CPU scheduler goes around the ready queue, allocating the CPU to each process for a time interval of up to 1 time quantum.



Implementation of Round Robin scheduling:

- We keep the ready queue as a FIFO queue of processes.
- New processes are added to the tail of the ready queue.
- The CPU scheduler picks the first process from the ready queue, sets a timer to interrupt after 1 time quantum, and dispatches the process.



One of two things will then happen

The process may have a CPU burst of less than 1 time quantum

The process itself will release the CPU voluntarily

The CPU scheduler will then proceed to the next process in the ready queue

The CPU burst of the currently running process is longer than 1 time quantum, the timer will go off and will cause an interrupt to the O.S.

A context switch will be executed, and the process will be put at the tail of the ready queue

The CPU scheduler will then select the next process in the ready queue



Round-Robin Scheduling (Turnaround Time and Waiting Time)

Consider the following set of processes that arrive at time 0, with the length of the CPU burst given in milliseconds and time quantum taken as 4 milliseconds for RR Scheduling:

Process ID	Burst Time
P1	24
P2	3
P3	3



Time Quantum

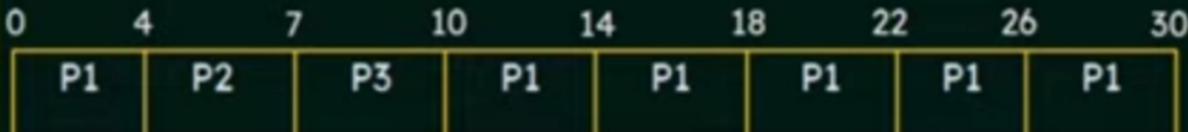
Gantt Chart:



Process ID	Burst Time
P1	24
P2	3
P3	3

Gantt Chart:

Time Quantum



Method 1

Turn Around time = Completion time - Arrival time

Waiting time = Turn Around time - Burst time

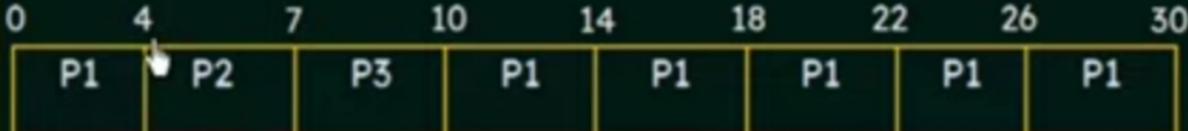
Process ID	Completion Time	Turnaround Time	Waiting Time
P1	30	$30 - 0 = 30$	$30 - 24 = 6$
P2	7	$7 - 0 = 7$	$7 - 3 = 4$
P3	10	$10 - 0 = 10$	$10 - 3 = 7$



Process ID	Burst Time
P1	24
P2	3
P3	3

Gantt Chart:

Time Quantum



Method 1

Turn Around time = Completion time - Arrival time

Waiting time = Turn Around time - Burst time

Process ID	Completion Time	Turnaround Time	Waiting Time
P1	30	$30 - 0 = 30$	$30 - 24 = 6$
P2	7	$7 - 0 = 7$	$7 - 3 = 4$
P3	10	$10 - 0 = 10$	$10 - 3 = 7$

Method 2

Waiting time = Last Start Time - Arrival Time -
(Preemption x Time Quantum)

Process ID	Waiting Time
P1	$26 - 0 - (5 \times 4) = 6$
P2	$4 - 0 - (0 \times 4) = 4$
P3	



Scheduling Algorithms

(Multilevel Queue Scheduling)

A class of scheduling algorithms has been created for situations in which processes are easily classified into different groups.

Example:

**Foreground
Processes
(Interactive)**

They have:

- Different response-time requirements
- Different scheduling needs

**Background
Processes
(Batch)**

In addition, foreground processes may have priority (externally defined) over background processes.

A multilevel queue scheduling algorithm partitions the ready queue into several separate queues.



A multilevel queue scheduling algorithm partitions the ready queue into several separate queues.

- The processes are permanently assigned to one queue, generally based on some property of the process, such as memory size, process priority, or process type.
- Each queue has its own scheduling algorithm.

Example:

Separate queues might be used for foreground and background processes.

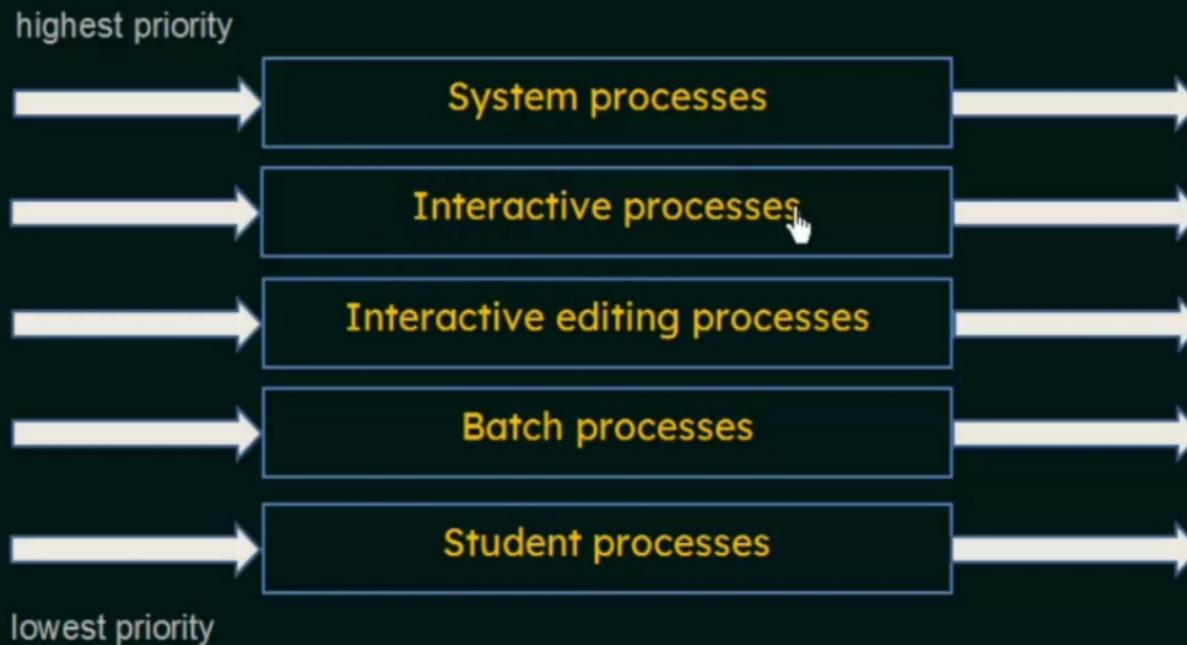
The foreground queue might be scheduled by an RR algorithm, while the background queue is scheduled by an FCFS algorithm.

In addition, there must be scheduling among the queues, which is commonly implemented as fixed-priority preemptive scheduling.

For example, the foreground queue may have absolute priority over the background queue. 



An example of a multilevel queue scheduling algorithm with five queues, listed below in order of priority:



Scheduling Algorithms

(Multilevel Feedback-Queue Scheduling)

The multilevel feedback-queue scheduling algorithm allows a process to move between queues.

- The idea is to separate processes according to the characteristics of their CPU bursts.
- If a process uses too much CPU time, it will be moved to a lower-priority queue.
- This scheme leaves I/O-bound and interactive processes in the higher-priority queues.
- In addition, a process that waits too long in a lower-priority queue may be moved to a higher-priority queue.

This form of aging prevents starvation.





Figure: Multilevel feedback queues



In general, a multilevel feedback-queue scheduler is defined by the following parameters:

- The number of queues
- The scheduling algorithm for each queue
- The method used to determine when to upgrade a process to a higher priority queue
- The method used to determine when to demote a process to a lower priority queue
- The method used to determine which queue a process will enter when that process needs service

